

DATABASE MANAGEMENT SYSTEM

Chapter 10 Data Evaluation and Transaction

Bikash Khadka Shah

MCSE, MSDA, OCP, RHCE, RHCVA, CCNA, CEH

9841766620 | 9801076620



Transactions

Definition

- A transaction is a collection of one or more operations on one or more databases, which reflects a single real-world transition
 - In the real world, this happened (completely) or it didn't happen at all (**Atomicity**)
- Commerce examples
 - Transfer money between accounts
 - Purchase a group of products
- Student record system
 - Register for a class (either waitlist or allocated)

Coding a transaction

- Typically a computer-based system doing OLTP has a collection of *application programs*
- Each program is written in a high-level language, which calls DBMS to perform individual SQL statements
 - Either through embedded SQL converted by preprocessor
 - Or through Call Level Interface where application constructs appropriate string and passes it to DBMS

Why write programs?

- Why not just write a SQL statement to express “what you want”?
- An individual SQL statement can't do enough
 - It can't update multiple tables
 - It can't perform complicated logic (conditionals, looping, etc)

COMMIT

- As app program is executing, it is “in a transaction”
- Program can execute COMMIT
 - SQL command to finish the transaction successfully
 - The next SQL statement will automatically start a new transaction

Warning

- The idea of a transaction is hard to see when interacting directly with DBMS, instead of from an app program
- Using an interactive query interface to DBMS, by default each SQL statement is treated as a separate transaction (with implicit COMMIT at end) unless you explicitly say “START TRANSACTION”

A Limitation

- Some systems rule out having both DML and DDL statements in a single transaction
- i.e., you can change the schema, or change the data, but not both

ROLLBACK

- If the app gets to a place where it can't complete the transaction successfully, it can execute ROLLBACK
- This causes the system to “abort” the transaction
 - The database returns to the state without any of the previous changes made by activity of the transaction

Reasons for Rollback

- User changes their mind (“ctl-C”/cancel)
- Explicit in program, when app program finds a problem
 - e.g. when qty on hand < qty being sold
- System-initiated abort
 - System crash
 - Housekeeping
 - e.g. due to timeouts

Atomicity

- Two possible outcomes for a transaction
 - It *commits*: all the changes are made
 - It *aborts*: no changes are made
- That is, transaction's activities are **all** or **nothing**

Integrity

- A real world state is reflected by collections of values in the tables of the DBMS
- But not every collection of values in a table makes sense in the real world
- The state of the tables is restricted by **integrity constraints**
- e.g. account number is unique
- e.g. stock amount can't be negative

Integrity

- Many constraints are explicitly declared in the schema
 - So the DBMS will enforce them
 - Especially: primary key (some column's values are non null, and different in every row)
 - And referential integrity: value of foreign key column is actually found in another "referenced" table
- Some constraints are not declared
 - They are business rules that are supposed to hold

Consistency

- Each transaction can be written on the assumption that all integrity constraints hold in the data, before the transaction runs
- It must make sure that its changes leave the integrity constraints still holding
 - However, there are allowed to be intermediate states where the constraints do not hold
- A transaction that does this, is called **consistent**
- This is an obligation on the programmer
 - Usually the organization has a testing/checking and sign-off mechanism before an application program is allowed to get installed in the production system

System obligations

- Provided the app programs have been written properly,
- Then the DBMS is supposed to make sure that the state of the data in the DBMS reflects the real world accurately, as affected by all the committed transactions

Local to global reasoning

- Organization checks each app program as a separate task
 - Each app program running on its own moves from state where integrity constraints are valid to another state where they are valid
- System makes sure there are no nasty interactions
- So the final state of the data will satisfy all the integrity constraints

Example - Tables

- System for managing inventory
- InStore(prodID, storeID, qty)
- Product(prodID, desc, mnfr, ..., WarehouseQty)
- Order(orderNo, prodID, qty, rcvd,)
 - Rows never deleted!
 - Until goods received, rcvd is null
- Also Store, Staff, etc etc

Example - Constraints

- **Primary keys**
 - **InStore: (prodID, storeID)**
 - **Product: prodID**
 - **Order: orderId**
 - **etc**
- **Foreign keys**
 - **Instore.prodID references Product.prodID**
 - **etc**

Example - Constraints

- Data values
 - `Instore.qty` ≥ 0
 - `Order.rcvd` \leq `current_date` or `Order.rcvd` is null
- Business rules
 - for each p, (Sum of qty for product p among all stores and warehouse) ≥ 50
 - for each p, (Sum of qty for product p among all stores and warehouse) ≥ 70 or there is an outstanding order of product p

Example - transactions

- **MakeSale(store, product, qty)**
- **AcceptReturn(store, product, qty)**
- **RcvOrder(order)**
- **Restock(store, product, qty)**
 - **// move from warehouse to store**
- **ClearOut(store, product)**
 - **// move all held from store to warehouse**
- **Transfer(from, to, product, qty)**
 - **// move goods between stores**

Example - ClearOut

- Validate Input (appropriate product, store)
- SELECT qty INTO :tmp
FROM InStore
WHERE StoreID = :store AND prodID = :product
- UPDATE Product
SET WarehouseQty = WarehouseQty + :tmp
WHERE prodID = :product
- UPDATE InStore
SET Qty = 0
WHERE prodID = :product
- COMMIT

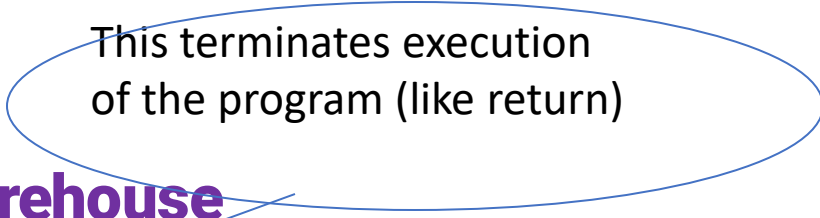
This is one way to write the application; other algorithms are also possible

Example - Restock

- Input validation
 - Valid product, store, qty
 - Amount of product in warehouse \geq qty
- UPDATE Product
 - SET WarehouseQty = WarehouseQty - :qty
 - WHERE prodID = :product
- If no record yet for product in store
 - INSERT INTO InStore (:product, :store, :qty)
- Else, UPDATE InStore
 - SET qty = qty + :qty
 - WHERE prodID = :product and storeID = :store
- COMMIT

Example - Consistency

- How to write the app to keep integrity holding?
- MakeSale logic:
 - Reduce Instore.qty
 - Calculate sum over all stores and warehouse
 - If sum < 50, then ROLLBACK // Sale fails
 - If sum < 70, check for order where date is null
 - If none found, insert new order for say 25
 - COMMIT



This terminates execution of the program (like return)

Example - Consistency

- We don't need any fancy logic for checking the business rules in Restock, ClearOut, Transfer
 - Because sum of qty not changed; presence of order not changed
 - *provided integrity holds before txn, it will still hold afterwards*
- We don't need fancy logic to check business rules in AcceptReturn
 - *why?*
- Is checking logic needed for RcvOrder?

Threats to data integrity

- Need for application rollback
- System crash
- Concurrent activity
- The system has mechanisms to handle these

Application rollback

- A transaction may have made changes to the data before discovering that these aren't appropriate
 - the data is in state where integrity constraints are false
 - Application executes ROLLBACK
- System must somehow return to earlier state
 - Where integrity constraints hold
- So aborted transaction has no effect at all

Example

- While running MakeSale, app changes InStore to reduce qty, then checks new sum
- If the new sum is below 50, txn aborts
- System must change InStore to restore previous value of qty
 - Somewhere, system must remember what the previous value was!

System crash

- At time of crash, an application program may be part-way through (and the data may not meet integrity constraints)
- Also, buffering can cause problems
 - Note that system crash loses all buffered data, restart has only disk state
 - Effects of a committed txn may be only in buffer, not yet recorded in disk state
 - Lack of coordination between flushes of different buffered pages, so even if current state satisfies constraints, the disk state may not

Example

- Suppose crash occurs after
 - MakeSale has reduced InStore.qty
 - found that new sum is 65
 - found there is no unfilled order
 - // but before it has inserted new order
- At time of crash, integrity constraint did not hold
- Restart process must clean this up (effectively aborting the txn that was in progress when the crash happened)

Concurrency

- When operations of concurrent threads are interleaved, the effect on shared state can be unexpected
- Well known issue in operating systems, thread programming
 - see OS textbooks on critical section
 - Java use of synchronized keyword

Famous anomalies

- **Dirty data**
 - One task T reads data written by T' while T' is running, then T' aborts (so its data was not appropriate)
- **Lost update**
 - Two tasks T and T' both modify the same data
 - T and T' both commit
 - Final state shows effects of only T, but not of T'
- **Inconsistent read**
 - One task T sees some but not all changes made by T'
 - The values observed may not satisfy integrity constraints
 - This was not considered by the programmer, so code moves into absurd path

Example – Dirty data

AcceptReturn(p1,s1,50)

MakeSale(p1,s2,65)

Update row 1: 25 -> 75

update row 2: 70->5

find sum: 90

// no need to insert

// row in Order

Abort

// rollback row 1 to 25

Integrity constraint is false:
Sum for p1 is only 40!

COMMIT

p1	s1	25
p1	s2	70
p2	s1	60
etc	etc	etc

p1	etc	10
p2	etc	44
etc	etc	etc

Initial state of InStore, Product

p1	s1	25
p1	s2	5
p2	s1	60
etc	etc	etc

p1	etc	10
p2	etc	44
etc	etc	etc

Final state of InStore, Product

Example – Lost update

- **ClearOut(p1,s1)**
AcceptReturn(p1,s1,60)
- **Query InStore; qty is 25**
- **Add 25 to WarehouseQty: 40->65**
- **Update row 1: 25->85**
- **Update row 1, setting it to 0**
- **COMMIT**
- **COMMIT**

60 returned p1's have vanished from system;
total is still 135

p1	s1	25
p1	s2	50
p2	s1	45
etc	etc	etc

p1	etc	40
p2	etc	55
etc	etc	etc

Initial state of InStore, Product

p1	s1	0
p1	s2	50
p2	s1	45
etc	etc	etc

p1	etc	65
p2	etc	55
etc	etc	etc

Final state of InStore, Product

Example – Inconsistent read

- ClearOut(p1,s1) **MakeSale(p1,s2,60)**
- Query InStore: qty is 30
- Add 30 to WarehouseQty: 10->40
- **update row 2: 65->5**
- **find sum: 75**
- **// no need to**
- **insert**
- **// row in Order**
- Update row 1, setting it to 0
- COMMIT
- **COMMIT**

Integrity constraint is false:
Sum for p1 is only 45!

p1	s1	30
p1	s2	65
p2	s1	60
etc	etc	etc

Initial state of InStore, Product

p1	etc	10
p2	etc	44
etc	etc	etc

p1	s1	0
p1	s2	5
p2	s1	60
etc	etc	etc

p1	etc	40
p2	etc	44
etc	etc	etc

Final state of
InStore, Product

Serializability

- To make isolation precise, we say that an execution is serializable when
- There exists some serial (ie batch, no overlap at all) execution of the same transactions which has the same final state
 - Hopefully, the real execution runs faster than the serial one!
- NB: different serial txn orders may behave differently; we ask that *some* serial order produces the given state
 - Other serial orders may give different final states

Example – Serializable execution

- **ClearOut(p1,s1)**
- **Query InStore: qty is 30**
- **MakeSale(p1,s2,20)**
- **update row 2: 45->25**
- **find sum: 65**
- **no order for p1 yet**
- **Add 30 to WarehouseQty: 10->40**
- **Update row 1, setting it to 0**
- **COMMIT**
- **Insert order for p1**
- **COMMIT**

Execution is like serial
MakeSale; ClearOut

p1	s1	30
p1	s2	45
p2	s1	60
etc	etc	etc

p1	etc	10
p2	etc	44
etc	etc	etc

Order: empty

Initial state of InStore, Product, Order

p1	s1	0
p1	s2	25
p2	s1	60
etc	etc	etc

p1	etc	40
p2	etc	44
etc	etc	etc

p1	25	Null	etc
----	----	------	-----

Final state of InStore, Product, Order

Serializability Theory

- There is a beautiful mathematical theory, based on formal languages
 - Treat the set of all serializable executions as an object of interest (called **SR**)
 - Thm: SR is in NP, i.e. the task of testing whether an execution is serializable seems unreasonably slow
- Does it matter?
 - The goal of practical importance is to design a system that produces some subset of the collection of serializable executions
 - It's not clear that we care about testing arbitrary executions that don't arise in our system

Conflict serializability

- There is a nice sufficient condition (ie a conservative approximation) called **conflict serializable**, which can be efficiently tested
 - Draw a **precedes graph** whose nodes are the transactions
 - Edge from T_i to T_j when T_i accesses x , then later T_j accesses x , and the accesses conflict (not both reads)
 - The execution is conflict serializable iff the graph is acyclic
- Thm: if an execution is conflict serializable then it is serializable
 - Pf: the serial order with same final state is any topological sort of the precedes graph
- Most people and books use the approximation, usually without mentioning it!

ACID

- **A**tomic
 - State shows either all the effects of txn, or none of them
- **C**onsistent
 - Txn moves from a state where integrity holds, to another where integrity holds
- **I**solated
 - Effect of txns is the same as txns running one after another (ie looks like batch mode)
- **D**urable
 - Once a txn has committed, its effects remain in the database

Transaction: Atomicity

- The atomicity property of a transaction requires that all operations of a transaction be completed, if not, the transaction is aborted.
- In other words, a transaction is treated as single, individual logical unit of work.
- Therefore, a transaction must execute and complete each operation in its logic before it commits its changes.
- As the transaction is considered as one operation even though there are multiple read and writes it completes or fails as one unit.
- The atomicity property of transaction is ensured by the transaction recovery subsystem of a DBMS.
- In the event of a system crash in the midst of transaction execution, the recovery techniques undo any effects of the transaction on the database.

Transaction: Consistency

- **Execution of a transaction must leave a database in either its prior stable state or a new stable state that reflects the new modifications (updates) made by the transaction.**
 - If the transaction fails, the database must be returned to the state it was in prior to the execution of the failed transaction.
 - If the transaction commits, the database must reflect the new changes.
- **Thus, all resources are always in a consistent state.**
- **In other words, a transaction must transform the database from one consistent state to another consistent state.**
- **The preservation of consistency is generally the responsibility of the programmers who write the database programs or of the DBMS module that enforces integrity constraints.**

Transaction: Isolation

- Isolation property of a transaction means that the data used during the execution of a transaction cannot be used by a second transaction until the first one is completed.
- This property isolates transactions from one another.
- In other words, if a transaction T1 is being executed and is using the data item X, that data item cannot be accessed by any other transaction (T2.....Tn) until T1 ends.
- The isolation property is enforced by the concurrency control subsystem of the DBMS.

Transaction: Durability

- **The durability property of transaction indicates the performance of the database's consistent state.**
- **It states that the changes made by a transaction are permanent.**
- **They cannot be lost by either a system failure or by the erroneous operation of a faulty transaction.**
- **When a transaction is completed, the database reaches a consistent state and that state cannot be lost, even in the event of system's failure.**
- **Durability property is the responsibility of the recovery subsystem of the DBMS.**

Single User vs Multiuser System

- **A DBMS is single-user if at most one user at a time can use the system.**
- **It is multi user if many users can use the system at a time.**
 - **So they can access the database concurrently.**
- **Single user DBMSs are mostly restricted to a PC; most of the DBMSs are multiuser and many of the users would be accessing them for business process continuation.**

Concurrency control

- **Concurrency control and recovery mechanisms are mainly concerned with the database commands in the transaction.**
- **Transactions submitted by the various users may execute concurrently and may access and update the same database items.**
- **If the concurrent execution is uncontrolled, it may lead to problems:**
 - **The lost update problem**
 - **The temporary update (or dirty read) problem**
 - **The incorrect summary problem**
 - **The unrepeatable read problem**

Concurrency control...

- **Lost update problem occurs when two transactions that access the same database items have interleaved in a way that makes the value of some db item incorrect.**
- **Dirty read problem occurs when one transaction updates a database item and fails for some reason. Meanwhile, the updated db item is read by another transaction before it is changed back to its original value.**
- **If one transaction is calculating an aggregate summary function on a number of db items while other transactions are updating some of these items, the incorrect summary problem occurs.**
- **In unrepeatable read, a transaction T reads the same item twice and the item is changed by another transaction T1 between the two reads.**

Concurrency Control Methods

- **Concurrency control is provided in a database to:**
 - **enforce isolation among transactions.**
 - **preserve database consistency through consistency preserving execution of transactions.**
 - **resolve read-write and write-read conflicts.**
- **Concurrency control protocols can be broadly divided into two categories –**
 - **Lock based protocols**
 - **Time stamp based protocols**

Lock Based Protocols

- A lock is a variable associated with a data item that describes a status of data item with respect to possible operation that can be applied to it.
- They synchronize the access by concurrent transactions to the database items.
- In lock based protocols, It is required that all the data items must be accessed in a mutually exclusive manner.
- Shared Lock (S):
 - Also known as Read-only lock.
 - As the name suggests it can be shared between transactions because while holding this lock the transaction does not have the permission to update data on the data item.
- Exclusive Lock (X):
 - Data item can be both read as well as written.
 - This is Exclusive and cannot be held simultaneously on the same data item.
 - X-lock is requested using lock-X instruction.

Lock compatibility matrix:

	S	X
S	✓	✗
X	✗	✗

- **Lock conversion:**
 - A transaction that holds a lock on an item A is allowed under certain condition to change the lock state from one state to another.
 - **Upgrade:** A S(A) can be upgraded to X(A) if Ti is the only transaction holding the S-lock on element A.
 - **Downgrade:** We may downgrade X(A) to S(A) when we feel that we no longer want to write on data-item A. As we were holding X-lock on A, we need not check any conditions.

Lock Based Protocols...

- Simple locking may not always produce Serializable results, it may lead to Deadlock.
 - It allows transactions to obtain a lock on every object before beginning operation.
 - Transactions may unlock the data item after finishing the 'write' operation.
- However these may cause some problems, eg:

	T1	T2
1	lock-X(B)	
2	read(B)	
3	B:=B-50	
4	write(B)	
5		lock-S(A)
6		read(A)
7		lock-S(B)
8	lock-X(A)	
9

Lock Based Protocols...

- **Deadlock:**
 - Consider the above execution phase.
 - Now, T1 holds an Exclusive lock over B, and T2 holds a Shared lock over A.
 - Consider Statement 7, T2 requests for lock on B, while in Statement 8 T1 requests lock on A.
 - This as you may notice imposes a Deadlock as none can proceed with their execution.
- **Starvation:**
 - It is also possible if concurrency control manager is badly designed.
 - For example: A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.

Two Phase Locking

- A transaction is said to follow Two Phase Locking protocol if Locking and Unlocking can be done in two phases.
 - Growing Phase: New locks on data items may be acquired but none can be released.
 - Shrinking Phase: Existing locks may be released but no new locks can be acquired.
- If lock conversion is allowed, then upgrading of lock(from S(a) to X(a)) is allowed in Growing Phase and downgrading of lock (from X(a) to S(a)) must be done in shrinking phase.
- The Point at which the growing phase ends, i.e., when transaction takes the final lock it needs to carry on its work is known as Lock Point.
- 2PL ensures serializability, but cascading rollbacks and deadlocks are possible.

Two Phase Locking...

	T1	T2
1	LOCK-S(A)	
2		LOCK-S(A)
3	LOCK-X(B)	
4
5	UNLOCK(A)	
6		LOCK-X(C)
7	UNLOCK(B)	
8		UNLOCK(A)
9		UNLOCK(C)
10

- A transaction implementing 2-PL.
- Transaction T1:
 - Growing Phase is from steps 1-3, Shrinking Phase is from steps 5-7.
 - Lock Point at 3
- Transaction T2:
 - Growing Phase is from steps 2-6, Shrinking Phase is from steps 8-9.
 - Lock Point at 6

Two Phase Locking

- There are some variations of 2PL.
- **Strict 2-PL** requires that in addition to the lock being 2-Phase all Exclusive(X) Locks held by the transaction be released until after the Transaction Commits; ensuring cascade less and recoverable.
- **Rigorous 2-PL** requires that in addition to the lock being 2Phase all Exclusive(X) and Shared(S) Locks held by the transaction be released until after the Transaction Commits.
- **Conservative 2-PL or Static 2-PL**, requires the transaction to lock all the items it access before the Transaction begins execution by pre-declaring its read-set and write-set.
 - If any of the pre-declared items needed cannot be locked, the transaction does not lock any of the items, instead it waits until all the items are available for locking.

Evaluation of DBMS

To choose the most suitable DBMS, we need to evaluate from various systems. We perform a structured approach to evaluate the database systems.

Evaluation Methodology

1. Feature analysis

In here, we determine whether the DBMS provides all the features required for the operations that are to be performed on the data and shortlist the DBMS. This analysis serves two purposes:

- It directly eliminates the DBMS that is not competent to fulfill all the needs of the user.
- We can rank the remaining systems based on the features provided by them.

2. Performance Analysis

In here, we analyze only those shortlisted DBMS, evaluate the systems' efficiency, and choose one with maximum efficiency. The purpose of this analysis is:

- To analyze the behavior of the database management system in the user system.
- To gather the data on the performance of the DBMS and rank them accordingly.

Database evaluation process

1. Planning the Database.
2. Analyzing Information for the Database.
3. Designing the Database.
4. Implementing the Database.
5. Testing the Database.
6. Maintaining the Database.

Planning the Database

- **Gather facts to understand the problem the database is trying to solve**
- **Specify the overall objectives of the database**
- **Deliverable: Database requirements statements**

Analyzing Information for the Database

- **Identify functional and user requirements (what will be stored, how it will be used, expected reports)**
- **Deliverable: Database system specifications**

Designing the Database

- **Design each technical component based on the system specifications**
- **Document in a technical specifications document**
- **Deliverable: Conceptual database technical specifications**

Implementing the Database

- **Create the actual database design and install it**
- **May be phased to test each component**
- **Includes developing applications, testing, documentation, user training, populating data**
- **Deliverable: The operational database**

Testing the Database

- Test the database to ensure it meets all requirements
- Install on all required computers and test links between components
- Deliverable: Functional database ready for use

Maintaining the Database

- Monitor database functioning to ensure continuous operation
- Proactive, preventive, corrective and adaptive maintenance tasks
- Allocate user access, routine monitoring, periodic testing
- Deliverable: Ongoing database with backed-up data that evolves to meet changing requirements